

CrashCourse::Dplyr

Tinashe M. Tapera

1 / 28

Agenda

1. The Philosophy
2. The Basics
3. The Extras

The Philosophy

"**dplyr** is a grammar of data manipulation, providing a consistent set of **verbs** that help you solve the most common data manipulation challenges."

— <https://dplyr.tidyverse.org/>

- If R is a language, **dplyr** is a dialect
- Main focus on data munging within the R ecosystem:
 - You're gonna wanna use `tibble()`s
- Focuses on elegance, readability, parsimony, and reproducibility
- Part of the **tidyverse**, so works well with all of their packages



The Philosophy

- `dplyr` abstracts base R; does not replace direct knowledge
- Not very widely scoped; wouldn't use it for `out-of-tibble()` situations (but many situations in R can be manipulated into `tibble()`-friendly ones)
- Elegant, but not the fastest; with very large datasets, `data.table` is faster ([source](#))

The Basics

"I claim that most single table problems can be solved with just five key verbs: filter, select, mutate, arrange and summarise, along with a 'by group' adverb."

— Hadley Wickham

The Basics | %>%

Pipes conjoin each dplyr verb by saying "and then...".

```
library(dplyr, warn.conflicts = FALSE, quietly = TRUE)
```

```
## Warning: package 'dplyr' was built under R version 3.5.1
```

```
iris %>%  
  head() #note the indentation for readability
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species  
## 1           5.1         3.5          1.4          0.2   setosa  
## 2           4.9         3.0          1.4          0.2   setosa  
## 3           4.7         3.2          1.3          0.2   setosa  
## 4           4.6         3.1          1.5          0.2   setosa  
## 5           5.0         3.6          1.4          0.2   setosa  
## 6           5.4         3.9          1.7          0.4   setosa
```

The Basics | %>%

- Typically, you could just do `head(iris)`: parsimonious & readable!
- What if you had 3, 4, or more functions wrapping around one object?

e.g. what are the coefficients of a linear model predicting Species in iris, using only rows where Sepal.Length is greater than the mean Sepal Length?

```
coef(lm(Species ~., data=subset(iris, iris$Sepal.Length > mean(iris$Sepal.Length)))
```

```
## (Intercept) Sepal.Length Sepal.Width Petal.Length Petal.Width
## 0.4760524 -0.2514633 -0.2395932 0.4583072 0.6171173
```

- Some programmers suggest breaking up their compound lines by assigning outputs to variables incrementally, e.g.

```
meanSepal = mean(iris$Sepal.Length)
subsdf = subset(iris, iris$Sepal.Length > meanSepal)
# etc...
```

- Really, tho...?

The Basics | %>%

In dplyr, it looks like this:

```
iris%>%  
  filter(Sepal.Length > mean(Sepal.Length))%>%  
  lm(Species ~ ., data=.)%>%  
  coef()
```

the noun
first verb
second verb
last verb

```
## (Intercept) Sepal.Length Sepal.Width Petal.Length Petal.Width  
## 0.4760524 -0.2514633 -0.2395932 0.4583072 0.6171173
```


The Basics | .

- Formulas in R can make use of a period or dot operator: `lm(Species ~ ., data=iris)`. The dot refers to "all variables except those on the LHS/RHS".
- In `dplyr` (really, `magrittr`), the dot refers to the noun being passed around, and is implicit by default.
- It is described as a "dummy parameter" or "placeholder"

```
iris%>%                               # iris%>%
  head(2)                             #   head(x=., n=2)
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
## 1           5.1           3.5           1.4           0.2   setosa
## 2           4.9           3.0           1.4           0.2   setosa
```

- There are some nuances to its use that will come up as you get into more advanced operations

The Basics | `select()`

- Most important verb to understand
- For selecting columns or variables (as long as object is a kind of dataframe)

```
iris%>%  
  select(Sepal.Length, Petal.Length)%>%  
  head(3)
```

```
##   Sepal.Length Petal.Length  
## 1           5.1           1.4  
## 2           4.9           1.4  
## 3           4.7           1.3
```

- Compare with bracket `[,]` indexing:

```
head( iris[, grep("Length$", names(iris))], 3)
```

```
##   Sepal.Length Petal.Length  
## 1           5.1           1.4  
## 2           4.9           1.4  
## 3           4.7           1.3
```

The Basics | `select()`

So many helper functions with `select()`!!!

- – to drop, `Var1:Var5` for a range, `c(...)` for vectors
- `starts_with()`, `ends_with()`, `contains()`, `matches("regular_expression")` for regular expressions
- `one_of(c(...))` for optional matching
- `everything()` for everything that's left

```
iris %>%  
  select(ends_with("Length")) %>%  
  head(3)
```

```
##   Sepal.Length Petal.Length  
## 1           5.1           1.4  
## 2           4.9           1.4  
## 3           4.7           1.3
```

The Basics | `filter()`

- For selecting rows of a dataframe

```
iris%>%  
  filter(Species == "setosa")%>%  
  head(3)
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
## 1	5.1	3.5	1.4	0.2	setosa
## 2	4.9	3.0	1.4	0.2	setosa
## 3	4.7	3.2	1.3	0.2	setosa

- Compare with `which()`

```
head( iris[which(iris$Species == "setosa"), ] ,3)
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
## 1	5.1	3.5	1.4	0.2	setosa
## 2	4.9	3.0	1.4	0.2	setosa
## 3	4.7	3.2	1.3	0.2	setosa

The Basics | `mutate()`

- Use `mutate()` to add new columns onto the dataframe (and use `transmute()` to only return the new column)
- Implicitly calls `select()`

```
iris %>%  
  mutate(Sepal.Area = Sepal.Length*Sepal.Width) %>%  
  head(3)
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species	Sepal.Area
## 1	5.1	3.5	1.4	0.2	setosa	17.85
## 2	4.9	3.0	1.4	0.2	setosa	14.70
## 3	4.7	3.2	1.3	0.2	setosa	15.04

The Basics | `arrange()`

- Sort by variable(s)

```
iris %>%  
  arrange(-Sepal.Length, -Sepal.Width) %>%  
  head(3)
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
## 1	7.9	3.8	6.4	2.0	virginica
## 2	7.7	3.8	6.7	2.2	virginica
## 3	7.7	3.0	6.1	2.3	virginica

- Also respects grouping variables, too!

The Basics | summarise()

- Takes a range of rows and applies some function to return a dataframe of one value:

```
iris %>%  
  summarise(my_mean=mean(Petal.Length))
```

```
##      my_mean  
## 1      3.758
```

The Basics | `group_by()`

- Applies a grouping arrangement to be passed on to other functions further on
- All basic functions in `dplyr` respect grouping

```
iris %>%  
  group_by(Species) %>%  
  summarise(my_mean=mean(Petal.Length))
```

```
## # A tibble: 3 x 2  
##   Species      my_mean  
##   <fct>      <dbl>  
## 1 setosa      1.46  
## 2 versicolor  4.26  
## 3 virginica   5.55
```

- You can group by as many variables as you'd like

The Basics | `gather()`

- `gather()` is great for turning wide-form data into long-form. Args as follows:

```
gather(key = name the variable stack,
       value = name the values,
       ... = which variables to gather (use select() helpers))
```

```
library(tidyr)
iris %>%
  gather(Petal_metric, value, starts_with("Petal")) %>%
  head(10)
```

##		Sepal.Length	Sepal.Width	Species	Petal_metric	value
##	1	5.1	3.5	setosa	Petal.Length	1.4
##	2	4.9	3.0	setosa	Petal.Length	1.4
##	3	4.7	3.2	setosa	Petal.Length	1.3
##	4	4.6	3.1	setosa	Petal.Length	1.5
##	5	5.0	3.6	setosa	Petal.Length	1.4
##	6	5.4	3.9	setosa	Petal.Length	1.7
##	7	4.6	3.4	setosa	Petal.Length	1.4
##	8	5.0	3.4	setosa	Petal.Length	1.5
##	9	4.4	2.9	setosa	Petal.Length	1.4
##	10	4.9	3.1	setosa	Petal.Length	1.5

The Basics | `spread()`

- The complement to `gather()`

`spread(key = the variable to unstack,
value = the variable with your stacked values)`

```
library(tibble)
iris %>%
  rownames_to_column("index") %>% # ?!?!
  gather(Petal_metric, value, starts_with("Petal")) %>%
  spread(Petal_metric, value) %>%
  arrange(as.numeric(index)) %>% # ?!?!
  head(5)
```

##	index	Sepal.Length	Sepal.Width	Species	Petal.Length	Petal.Width
## 1	1	5.1	3.5	setosa	1.4	0.2
## 2	2	4.9	3.0	setosa	1.4	0.2
## 3	3	4.7	3.2	setosa	1.3	0.2
## 4	4	4.6	3.1	setosa	1.5	0.2
## 5	5	5.0	3.6	setosa	1.4	0.2

- ?!?! One quirk: `spread()` needs specific row indices to unravel its values; throw in an explicit column index and order

The Basics | Other Common Functions

- `separate()/unite()`
 - Complementary string column "split" and "concatenate"
- `*_join()`
 - Traditional SQL-style joins (but with a nicer interface than `merge()`, `sqldf`, etc.)
- `sample_n()/sample_frac()`
 - Sampling rows of a dataframe (with or without replacement)
 - Much clearer than `iris[sample(nrow(iris), n),]`
- `slice()`
 - Positional row indexing
- Remember to `ungroup()` explicitly!
- Remember to use `rowwise()` to iterate (like calling `apply(MARGIN=2)`), because **R does not like iterating rows naturally**; it's vectorised!

The Extras | list-columns

- You can `nest()` dataframes and lists in `dplyr` to create list-columns

```
iris%>%  
  nest(-Species)%>%  
  as_tibble()                # for viewing on this slide
```

```
## # A tibble: 3 x 2  
##   Species      data  
##   <fct>      <list>  
## 1 setosa    <data.frame [50 x 4]>  
## 2 versicolor <data.frame [50 x 4]>  
## 3 virginica <data.frame [50 x 4]>
```

- You can then map operations on the list's objects with the `purrr` package (that's a topic for another day, though)

The Extras | list-columns

- Remember `summarise()`? It only works if the return of the summary gives you a single value vector. Using list-columns can help us override this
- What are the quantiles of `Sepal.Length` for each species of iris?

```
#base R, you'd have to call this three times for each species
quantile( iris[which(iris$Species == "setosa"), "Sepal.Length"] )
```

```
##      0%   25%   50%   75%  100%
##    4.3   4.8   5.0   5.2   5.8
```

```
#dplyr without list-columns returns an error

iris %>%
  group_by(Species) %>%
  summarise(quant=quantile(Sepal.Length))
```

```
## Error in summarise_impl(.data, dots): Column `quant` must be length 1 (a s
```

The Extras | list-columns

```
# instead, just coerce the return value into a list-column
iris %>%
  group_by(Species) %>%
  summarise(Sepal_Length_quants = list(quantile(Sepal.Length)))
```

```
## # A tibble: 3 x 2
##   Species      Sepal_Length_quants
##   <fct>        <list>
## 1 setosa      <dbl [5]>
## 2 versicolor <dbl [5]>
## 3 virginica  <dbl [5]>
```

The Extras | list-columns

```
iris %>%  
  group_by(Species) %>%  
  summarise(Sepal_Length_quants = list(quantile(Sepal.Length))) %>%  
  .$Sepal_Length_quants
```

```
## [[1]]  
##      0%   25%   50%   75%  100%  
##    4.3   4.8   5.0   5.2   5.8  
##  
## [[2]]  
##      0%   25%   50%   75%  100%  
##    4.9   5.6   5.9   6.3   7.0  
##  
## [[3]]  
##      0%   25%   50%   75%  100%  
##  4.900  6.225  6.500  6.900  7.900
```

The Extras | Scoped `filter_*()`

- You can create complex filtering conditions using scoped filters like `filter_at()`, `filter_all()`, and `filter_if()`
- These will return the rows of a dataframe once filtered on the specific variable predicates

e.g. let's filter only length variables from iris, where the length is greater than 5 for either of them

```
iris %>%  
  filter_at(.vars = vars(contains("Length")),  
            .vars_predicate = any_vars(. > 5)) %>%  
  head(3)
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
## 1	5.1	3.5	1.4	0.2	setosa
## 2	5.4	3.9	1.7	0.4	setosa
## 3	5.4	3.7	1.5	0.2	setosa

- Note the use of `vars()`, which explicitly calls `select()` helpers

The Extras | Scoped `mutate_*()`

- Similarly, there are scoped mutate and summarise calls

e.g. multiply only numeric variables by 2

```
iris %>%  
  transmute_if(is.numeric,  
               funs(new = . * 2)) %>%  
  head(3)
```

##	Sepal.Length_new	Sepal.Width_new	Petal.Length_new	Petal.Width_new
## 1	10.2	7.0	2.8	0.4
## 2	9.8	6.0	2.8	0.4
## 3	9.4	6.4	2.6	0.4

The Extras | Scoped `summarise_*()`

e.g. Let's summarise only width variables to get their means

```
# admittedly, this is overcomplicated
iris %>%
  summarise_at(vars(ends_with("width")),
               funs(mean)) %>%
  head(3)
```

```
##   Sepal.Width Petal.Width
## 1      3.057333      1.199333
```

- Note the use of `funs()`, which can accept any number of custom functions

Conclusion

- `dplyr` makes data munging cleaner and more interpretable
- There are lots of useful hidden functions under the hood
- Doesn't replace knowledge of base R
- Doesn't scale to HPC scenarios; best for making table summarisations and operations easier
- Huge community support means that you can figure out pretty much anything eventually

Thank you!